

# Enchanting PythonS

*to crunch data ...*

JuanJo Ciarlante

[@xjjo](#) [bit.ly/jjo-cv](http://bit.ly/jjo-cv)

#PyDayMDZ

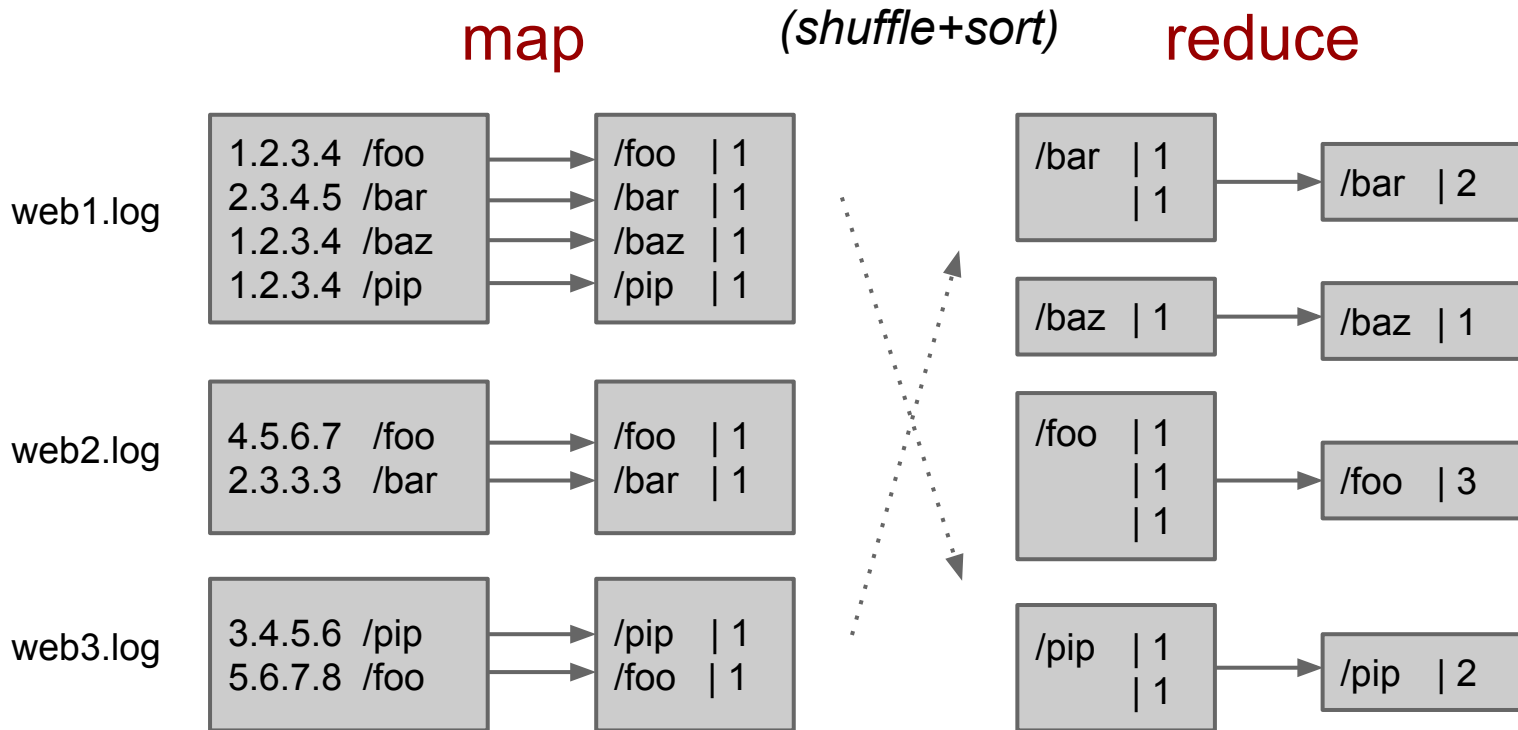
# MR: what ?

- framework for massive data processing
  - actually: data *transformation*
- based on 'rows'/records as:
  - <key,value>

# count'em all

- have: apache logs
- want: how many hits per page (urlpath) ?

# e.g.: hitcount by urlpath



# MR: how ?

- **map**: picks data from input rows
  - *record* ---> *key, data*
- (shuffle, sort) classifies by *key* to build:
  - *...* ---> *key, [data1, data2, ...]*
- **reduce**: aggregates, transforms - eg:
  - *key, [...]* ---> *key, sum([...])*

# MR: why is cool ?

- *kiss*:
  - really simple model
- *scalability*:
  - parallel-friendly by design
- *data-locality*:
  - distributed FS
- *sync-free*:
  - no explicit required IPC/sync between tasks

# **gimme that index**

- have: corpus of documents
- want: to search them by word (grep)

# e.g.: grep - filename by word

map

f1.txt

todos giran y  
giran ↵  
todos bajo el sol

todos		
giran		
<del>giran</del>		
todos		
bajo		
sol		?

f2.txt

me ha tomado el  
tiempo ↵ para  
verlos otra vez

me		
ha		
tomado		?
el		
...		

f3.txt

quién eeen ↵  
se ha tomado  
todo el vino

eeeen		
se		
ha		
tomado		?
todo		
vino		?

reduce

sol | ?

/idx/HH/sol.idx

tomado | ?  
| ?

/idx/HH/tomado.idx

vino | ?

/idx/HH/vino.idx

idx



# e.g.: grep - filename by word

map

f1.txt

todos giran y  
giran ↵  
todos bajo el sol

f2.txt

me ha tomado el  
tiempo ↵ para  
verlos otra vez

f3.txt

quién eeeen ↵  
se ha tomado  
todo el vino

todos |  
giran |  
~~giran~~ |  
todos |  
bajo |  
sol | f1.txt

me |  
ha |  
tomado | f2.txt  
el |  
... |

eeeen |  
se |  
ha |  
tomado | f3.txt  
todo |  
vino | f3.txt

reduce

sol | f1.txt

/idx/HH/sol.idx

tomado | f2.txt  
| f3.txt

/idx/HH/tomado.idx

vino | f3.txt

/idx/HH/vino.idx

idx

# MR: Hadoop

- floss \o/
- in Java :/, for Java :(
  - ¿ too much Javanic :-?
- => hadoop “streaming” \o/
  - arbitrary commands with pipelined data locality:  
*input* | **python mr.py**<sub>map</sub> | s+sort | **python mr.py**<sub>reduce</sub>

# MR: some python libs

- **MRJob**

- 👍 **local, hadoop**, Elastic MR (AWS)

- 👎 not hadoop 'native'

- **hadoopy**

- 👍 optimized for hadoop, **supports HDFS bin formats**

- 👎 only hadoop

- **discoproject.org**

- 👍 **100% python**

- 👎 python only, down to the DFS

# speaking of diversity ...

- have: apache logs
- want: to know how diversity of client IPs per page
  - shamelessly use `entropy(concatenated_IPs_bits)` as a *proxy* value for relative diversity


# e.g.: urlpath diversity

map


(shuffle+sort)

reduce: **entropy([ips])**









web1.log



web2.log



web3.log



2nd map reduce  
to aggregate url  
path by entropy



# MRjob: hitcount.py

<https://github.com/jjo/src-juanjo/blob/master/python/mrjob/j01-hitscount.py>

```
from mrjob.job import MRJob

class MRHitCount(MRJob):
    def mapper(self, _, line):
        ip, path = line.split()
        yield path, 1

    def reducer(self, key, values):
        yield key, sum(values)

if __name__ == '__main__':
    MRHitCount.run()
```

# MRjob: grep.py

<https://github.com/jjo/src-juanjo/blob/master/python/mrjob/j02-grep.py>

```
from mrjob.job import MRJob
from mrjob.compat import get_jobconf_value

class MRGrep(MRJob):
    def mapper(self, _, line):
        for word in line.split():
            yield word, get_jobconf_value('map.input.file')

    def reducer(self, key, values):
        yield key, str(values)

if __name__ == '__main__':
    MRGrep.run()
```

# MRjob: urlentropy.py

<https://github.com/jjo/src-juanjo/blob/master/python/mrjob/j04-entropy.py>

```
class MREntropyPerURL(MRJob):
    # 1st MR: urlpath -> entropy([ips])
    def input_mapper(self, _, line):
        ip, path = line.split()
        yield path, ip

    def urlpath_to_entropy(self, key, values):
        yield key, entropy_bits(values)

    # 2nd MR: aggregate all urlpaths by same entropy_val (omitted)
    # Pipe-line both MRs:
    def steps(self):
        return [self.mr(mapper=self.input_mapper, reducer=self.urlpath_to_entropy),
                self.mr(mapper=self.swap_values, reducer=self.values_per_key)]

if __name__ == '__main__':
    MREntropyPerURL.run()
```



# Hacktime \o/

- these slides:
  - <http://bit.ly/jjo-mrpy-14>
- some MR py libs:
  - mrjob, hadoop, happy
- interesting datasets:
  - <https://snap.stanford.edu/data/> networks
  - <http://aws.amazon.com/datasets/> diverse data
- complete source code for this slides
  - <https://github.com/jjo/src-juanjo/tree/master/python/mrjob>